

1 The Problem

In this report I programmed a gravity simulation that approximates the evolution of position and velocity of N stars in 2 dimensions. The code was written in C. The problem consists of reading a binary file containing information about position, velocity, mass and brightness of each star (the latter attribute was not used in this solution). The code was to be optimized for runtime while keeping the accuracy within a few magnitudes of order from the rounding errors of the datatypes used. This assignment was restricted to double types for the attributes of the stars. I used the modified force calculation called plummer-spheres, that reduces the risk of instability as stars get close and the denominator gets small by adding a constant epsilon:

$$\vec{F}_i = -Gm_i \sum_{j=0, j \neq i}^{N-1} \frac{m_j}{(r_{ij} + \epsilon_0)^3} \vec{r}_{ij} \quad (1)$$

I used the euler symplectic update of the stars velocity and position, because it preserves the total energy of the system.

2 Serial code version (A3)

2.1 The Solution

All following performance tests were run using the following parameters:

- $N = 3000$
- $\Delta t = 10^{-5}$
- $\epsilon_0 = 10^{-3}$
- $G = \frac{100}{N}$
- steps = 100

Where N is the number of stars, Δt is the time stepsize, ϵ_0 is the dampening factor in the force calculation, G is the gravitational constant and steps is the number of updates to iterate through. N and steps was initialized as const ints and Δt , ϵ_0 and G as const doubles for higher precision.

For all the performance tests I verified that the code solved the problem with at least an error less than 10^{-13} that was validated using reference files with the correct final state of the system. Graphics was used to get a better understanding of the bugs in the code, but once the code functioned as intended the graphics was turned off to allow for a faster runtime. All timings were calculated as a mean over 3 runs, using -O3 compiler optimization. I will briefly highlight the relevant parts of the codes that did not do best in testing, and end the report with the full contents of the code that performed best. The timings of all codes as well as a deeper discussion of optimization techniques used, considered and discarded can be found under section "Performance and Discussion". All codes presented used epsilon as a global constant, and all codes finalized the run by writing the results to a new .gal file.

2.1.1 First version

My first iteration of a solution consisted of a struct type as follows:

```
/* Define struct for a celestial body */
typedef struct {
    double xp; //X Position
    double yp; //Y Position
    double mass;
    double xv; //X Velocity
    double yv; //Y Velocity
    double brightness; //Keep, but not used in this simulation
} cbody_t;
```

Initialized as:

```
cbody_t stars[N];
```

And the function that updates the stars were initialized with following input arguments:

```
void update_stars(cbody_t stars[], double G, int N, double delta_t){
    for (int i = 0; i < N; i++) {
        /* Initialize variables used in this scope */

        /* More code... */
    }
}
```

```
    for (int j = 0; j < N; j++) {
        if (j == i){
            continue;
        }
        /* More code... */

        /* Acceleration contribution to star i from star j */
        double ax = fx/stars[i].mass;
        double ay = fy/stars[i].mass;

        /* Speed contribution to star i from star j */
        sum_xv += delta_t*ax;
        sum_yv += delta_t*ay;
    }
    /* New speed of star i */
    stars[i].xv += sum_xv;
    stars[i].yv += sum_yv;

    /* New position of star i */
    stars[i].xp += delta_t*stars[i].xv;
    stars[i].yp += delta_t*stars[i].yv;
}
}
```

2.1.2 Second version

My second iteration of a solution kept the same structure for the struct `cbody_t`, but initialized using dynamic allocation with pointers instead.

```
cbody_t *stars = (cbody_t *)malloc(N*sizeof(cbody_t));
```

Meaning that now the variable `stars` is a pointer to an array of `cbody_t` structs, the size of `N`. This pointer was passed to the `update_stars` function. I changed the function `update_stars` accordingly as well as added an optimization to the inner loop, further discussed in the next section.

```
void update_stars(cbody_t *stars, double G, int N, double dt){
    /* Initialize variables used in this scope */
```

```
cbody_t *star_ptr_i = stars;
cbody_t *star_ptr_j = stars + 1;

/* More code... */
for (int i = 0; i < N; i++) {
    /* Initialize variables used in this scope */

    /* More code... */

    /* Add contribution from i to all j as well as all j to i */
    for (int j = i+1; j < N; j++) {
        /* More code... */

        /*add this contribution to j */
        (star_ptr_j)->xv += vj*rix;
        (star_ptr_j)->yv += vj*riy;
        star_ptr_j++;
    }
    /* New speed of star i as consequence of all other stars*/
    (star_ptr_i)->xv += -Gdt*fx_sum/im;
    (star_ptr_i)->yv += -Gdt*fy_sum/im;
    (star_ptr_i)->xp += ((star_ptr_i)->xv)*dt;
    (star_ptr_i)->yp += ((star_ptr_i)->yv)*dt;
    star_ptr_i++;
    star_ptr_j = star_ptr_i + 1;
}
}
```

2.1.3 Third and final version

In the final code I changed the structure of the struct cbody_t:

```
/* Define struct for a celestial body */
typedef struct {
    double* xp; //X Position
    double* yp; //Y Position
    double* mass;
    double* xv; //X Velocity
```

```
    double* yv; //Y Velocity  
    double* brightness; //Keep, but not used in this simulation  
} cbody_t;
```

And initialized as:

```
cbody_t stars;  
stars.xp = malloc(N * sizeof(double));  
stars.yp = malloc(N * sizeof(double));  
stars.mass = malloc(N * sizeof(double));  
stars.xv = malloc(N * sizeof(double));  
stars.yv = malloc(N * sizeof(double));  
stars.brightness = malloc(N * sizeof(double));
```

meaning that I now pass a structure of arrays instead of an array of structures. I initialized pointers in the function `update_stars` that each pointed at one of the arrays holding a specific attribute of all the stars in the simulation.




Figure 1: Initial galaxy state, $N = 10000$




Figure 2: Galaxy state after a few hundred steps

2.2 Performance and Discussion

2.2.1 CPU and compiler used in this assignment

All codes were tested on a virtual linux machine (fredholm) provided by Uppsala University, here are the CPU specifications:

- Vendor ID: GenuineIntel
- Model name: Intel(R) Xeon(R) CPU E5520 @ 2.27GHz
- CPU family: 6
- Model: 26
- Thread(s) per core: 2

- Core(s) per socket: 4
- Socket(s): 2
- Stepping: 5
- CPU max MHz: 2267,0000
- CPU min MHz: 1600,0000
- BogoMIPS: 4533.50

Compiler information: 4:11.2.0-1ubuntu1 amd64 GNU C compiler

2.2.2 Results

I got the following runtimes for my code versions, using the "time" command in ubuntu: First version

Real	User	Sys
49.227	49.145s	0.072s

Second version

Real	User	Sys
22.330s	22.262s	0.060s

Third and final version

Real	User	Sys
18.186s	18.177s	0.000s

OBS: after presentation I got a run with real = 11.146s, user = 11,135s and results were still correct for the same input.

2.2.3 Complexity analysis

It was hypothesized that this algorithm has a complexity $\mathcal{O}(N^2)$. This means that runtime as a function of input N should scale as follows:

$$t(N) = c * N^2 \quad (2)$$

Where c is a constant that depends on factors outside of the algorithm, such as hardware. If the algorithm adheres to this complexity, then we would expect to see the following increase in time each time N is doubled:

$$\frac{t(2N)}{t(N)} = \frac{c * (2N)^2}{c * N^2} = 4 \quad (3)$$

meaning that time increases with a factor of 4 when N is doubled. When running the time command in ubuntu for $N = 1000, 2000, 4000$ and 8000 i got the following results:




Figure 3: Results from complexity analysis

thus we can confirm that the algorithm does in fact have a complexity of $\mathcal{O}(N^2)$.

2.2.4 Discussion

2.2.5 First version

The first code that performed the worst did so as a result of bad implementation of data structures and storing. Since function `update_stars` is passed the entire array of structs including the values that may not be necessary, such as brightness, this caused a lot of unnecessary memory usage which slowed the program down significantly.

2.2.6 Second version

In my second version I instead opted to send the structs as pointers which almost halved the runtime. There was still some problems with data locality, since each pointer pointed to an entire struct, and it also became intuitively complicated to handle the pointers, since each pointer had to be iterated through the array of pointers continuously in order to iterate to the next star. In this version I also changed the function `update_stars`, specifically the inner loop. Since the gravitational force between two objects is symmetrical, that means that we can use the force calculation from j to i to also update the contribution to change in velocity and position from i to j . This means that we never have to calculate the gravitational force between to objects more than once, and the number of iterations needed is greatly reduced.

2.2.7 Third and final version

The problem with the struct in the previous version was solved in the final version by converting the structure from an array of structs (AoS) to a struct of arrays (SoA). This meant that in the second version when operations of multiple attributes of the same struct was performed, the memory access pattern needs to "jump around" which can lead to poor cache locality and decrease performance. With a single struct, where each attribute is an array holding that specific attribute for all stars, all attributes of the same type lies in sequential order and there is no need for pointers to jump around in memory. Ultimately it did not give a significant improvement, but reduced

the runtime by about 2-3 seconds nonetheless. Furthermore, I initialized new variables from an expression wherever that expression was used more than once, so as to reduce the number of operations needed. I also initialized and assigned values to all variables that does not depend on i or j outside the loops, as well as initializing variables that only depends of i inside this loop, but outside the j loop to keep the inner loop as small as possible since this is usually good practice.

I considered using a hard rule that would stop calculations of forces between two stars if they were too far apart to have a significant force between them, although that may beat the purpose of using the double data type. I tried to analyze to see if any forces equalled zero from some rounding error, but with the double data type no forces was detected to be small enough for it not being able to be stored in a double type. I did not consider inlining since the only function in this code is relatively large with loops that would outweigh the performance decrease for that of the overhead of calling the function. I did not figure out any way to break the function down in to smaller chunks to inline the smaller functions instead. this code could be improved through parallelization, but this assignment was limited to serial codes only.

There was a compiler warning for my final code that is related to the fact that I do not check the return value for the `fread` function, which is good practice so the code can signal when the file was not able to be read. I did not change this although the performance increase for not adding this condition should be negligible, but it did not affect my results since it was confirmed using reference files that the code actually worked as intended.

It should also be added that the time measurements were done at different times of the day, and it seemed that I got better results when I ran the final code at around 04.00 (the time presented in this report) in the morning compared to around 13.00 (about +1 second in run time). This could perhaps be due to the system that was used is a shared distributed system where there are more active users during the day.

2.3 References

Initially I used the following code as starting point:

<https://github.com/LapuhRok/GalaxySimulation/blob/master/galsim.c>

although I ended up using very little from this code, apart from the function `void updateForceBasic()` which is very similar to my code, except that I applied certain optimizations and declared variables within their respective scope, as well as using pointers to update the stars velocity and position.

2.4 My final code

```
#include <stdio.h>
#include <stdlib.h>
#include "graphics.h"
#include <math.h>

/* Define struct for a celestial body */
typedef struct {
    double* xp; //X Position
    double* yp; //Y Position
    double* mass;
    double* xv; //X Velocity
    double* yv; //Y Velocity
    double* brightness; //Keep, but not used in this simulation
} cbody_t;

const float circleRadius=0.0025, circleColor=0;
const int windowHeight=800;
const double eps = 0.001;
//int check_counter = 0;

/* Calculates and updates for new velocity and position of single star */
void update_stars(cbody_t *stars, const double G, const int N, const double dt){
    double *xp = stars->xp;
    double *yp = stars->yp;
    double *mass = stars->mass;
    double *xv = stars->xv;
    double *yv = stars->yv;
```

```
const double Gdt = G*dt;
for (int i = 0; i < N; i++) {
    double im = mass[i];
    double ix = xp[i];
    double iy = yp[i];
    double axi = 0;
    double ayi = 0;
    for (int j = i+1; j < N; j++) {
        double mj = mass[j];
        double rx = ix - xp[j];
        double ry = iy - yp[j];
        double rsqrt = sqrt(rx*rx + ry*ry) + eps;
        double inv_r = 1.0/(rsqrt*rsqrt*rsqrt);
        double inv_r_mj = inv_r*mj;
        axi += rx*inv_r_mj;
        ayi += ry*inv_r_mj;
        xv[j] += Gdt*im*inv_r*rx;
        yv[j] += Gdt*im*inv_r*ry;
    }
    xv[i] += -Gdt*axi;
    yv[i] += -Gdt*ayi;
    xp[i] += xv[i]*dt;
    yp[i] += yv[i]*dt;
}

int main(int argc, char const *argv[]) {
    /* Make sure input arguments are correct */
    if (argc != 6) {
        printf("Error: Correct inputs: ./galsim N filename nsteps delta_t\n");
        return 1;
    }

    /* Save initial arguments and create the stars objects */
    const int N = atoi(argv[1]);
    const char *filename = argv[2];
    const int nsteps = atoi(argv[3]);
    const double dt = atof(argv[4]);
```

```
const int graphics = atoi(argv[5]);
int count = 0;
const double G = 100/(double)N;

//double G = 1.0;
printf("N = %d, filename = %s, nsteps = %d, delta_t = %lf, graphics = %d\n",
//CBody stars[N];
cbody_t stars;
stars.xp = malloc(N * sizeof(double));
stars.yp = malloc(N * sizeof(double));
stars.mass = malloc(N * sizeof(double));
stars.xv = malloc(N * sizeof(double));
stars.yv = malloc(N * sizeof(double));
stars.brightness = malloc(N * sizeof(double));

/* Read the file */
FILE *file = fopen(filename, "r");
if (file == NULL) {
    printf("Error: failed to open the file %s\n", filename);
    return 1;
}

/* Input initial conditions into stars */
/*for (i = 0; i < N; i++) {
    fread(&stars[i], sizeof(CBody), 1, file);
}*/

for (int i = 0; i < N; i++) {
    //cbody_t *new_star = &stars[i];
    //fread(, sizeof(cbody_t), 1, file);
    fread(&stars.xp[i], sizeof(double), 1, file);
    fread(&stars.yp[i], sizeof(double), 1, file);
    fread(&stars.mass[i], sizeof(double), 1, file);
    fread(&stars.xv[i], sizeof(double), 1, file);
    fread(&stars.yv[i], sizeof(double), 1, file);
    fread(&stars.brightness[i], sizeof(double), 1, file);
}
```

```
fclose(file); //Don't need file anymore

/* print values to check its working */
for (int i = 0; i < N; i++) {
    printf("Star #%d\n", i+1);
    printf("xp = %.15lf\n", stars.xp[i]);
    printf("yp = %.15lf\n", stars.yp[i]);
    printf("xv = %.15lf\n", stars.xv[i]);
    printf("yv = %.15lf\n", stars.yv[i]);
    printf("brightness = %.15lf\n", stars.brightness[i]);
    printf("mass = %.15lf\n", stars.mass[i]);
}

if (graphics == 1){
    printf("Graphics on\n");
    float L=1, W=1;

    InitializeGraphics(argv[0],windowWidth,windowWidth);
    SetCAxes(0,1);

    printf("Hit q to quit.\n");
    while(count < nsteps) {
        /* Update star position and velocity */
        update_stars(&stars, G, N, dt);
        /* Call graphics routines. */
        ClearScreen();
        for (int k = 0; k < N; k++) {
            DrawCircle(stars.xp[k], stars.yp[k], L, W, circleP
        }
        Refresh();
        /* Sleep a short while to avoid screen flickering. */
        usleep(3000);
        count++;
    }
    FlushDisplay();
    CloseDisplay();
} else {
    printf("Graphics turned off\n");
```

```
        /*for (i = 0; i < N; i++) {
            stars[i].xp += delta_t*stars[i].xv;
            stars[i].yp += delta_t*stars[i].yv;
        }*/
        while (count < nsteps) {
            update_stars(&stars, G, N, dt);
            count++;
        }/*
        for (i = 0; i < N; i++) {
            stars[i].xp += delta_t*stars[i].xv;
            stars[i].yp += delta_t*stars[i].yv;
        }*/
    }
    //printf("HERE COMES COUNT");
    //printf("Count = %d\n", count);
    /* Write result to new file */
    FILE *new_file = fopen("output.gal", "wb");
    //cbody_t *starPointer = stars;
    for (int i = 0; i < N; i++){
        fwrite(&stars.xp[i], sizeof(double), 1, new_file);
        fwrite(&stars.yp[i], sizeof(double), 1, new_file);
        fwrite(&stars.mass[i], sizeof(double), 1, new_file);
        fwrite(&stars.xv[i], sizeof(double), 1, new_file);
        fwrite(&stars.yv[i], sizeof(double), 1, new_file);
        fwrite(&stars.brightness[i], sizeof(double), 1, new_file);
    }
    fclose(new_file);
    free(stars.xp);
    free(stars.yp);
    free(stars.mass);
    free(stars.xv);
    free(stars.yv);
    free(stars.brightness);
    return 0;
}
```

3 Parallel code version (A4)

3.1 The solution

All parameters retain the same value as used and presented in The solution in part 2 of this report. Both codes were tested using the same reference data and was verified to get an accurate result before timing of the code proceeded. In this part the focus is solely on the performance and result, so no graphics was used and that part of the code is commented out.

3.2 Pthreads

The first change was to introduce a new struct type, which will hold all the variables that each thread gets a private copy of as such:

```
typedef struct {
    cbody_t *stars; // Pointer to celestial body
    int start; // Index where thread starts
    int end; // Index where thread ends
    double t_G; // Graviational constant
    double t_dt; //Time step
    int t_N; //N stars
} thread_info_t;
```

and initialized a global barrier, pthread_barrier_t barrier which will be used to synchronize threads where needed. The function update_stars gets the following update:

```
void* update_stars(void *arg){
    thread_info_t *tinfo = (thread_info_t *)arg;
    cbody_t *stars = tinfo->stars;
    const int start = tinfo->start;
    const int end = tinfo->end;
    const double t_G = tinfo->t_G;
    const double t_dt = tinfo->t_dt;
    const double Gdt = t_G*t_dt;
```



```

    const int N = tinfo->t_N;

    /* More code ... */
    for (i = start; i < end; i++) {
        /* More code .. */
        for (j = 0; j < N; j++) {
            jx = xp[j];
            jy = yp[j];
            mj = mass[j];
            rx = ix - jx;
            ry = iy - jy;
            rsqrt = sqrt(rx*rx + ry*ry) + eps;
            inv_r = 1.0/(rsqrt*rsqrt*rsqrt);
            inv_r_mj = inv_r*mj;
            axi += rx*inv_r_mj;
            ayi += ry*inv_r_mj;
        }
        xv[i] += -Gdt*axi;
        yv[i] += -Gdt*ayi;
    }
    pthread_barrier_wait(&barrier);
    for (int i = start; i < end; i++){
        xp[i] += xv[i]*t_dt;
        yp[i] += yv[i]*t_dt;
    }
    return NULL;
}

```

So now all previous input arguments are stored in a struct that the thread holds. The inner loop had to be downgraded to starting from $j = 0$ instead of from $j = i + 1$, since I no longer update any stars velocity in the inner loop since this solution resulted in race conditions. Race conditions also occurred for updating position, which had to be separated to a new loop since the updating of position depends on velocity. Therefore all threads finish to update all stars new velocity for the current timestep, before waiting at a barrier until all threads can continue with updating the position.

The following is added/modified to main():

```
int main(int argc, char const *argv[]) {

    /* More code ... */

    const int NUM_THREADS = atoi(argv[6]);
    pthread_t threads[NUM_THREADS];
    thread_info_t thread_info[NUM_THREADS];
    pthread_barrier_init(&barrier, NULL, NUM_THREADS);

    /* More code .. */

    while (count < nsteps){
        //update_stars(&stars, G, N, dt);
        for (int i = 0; i < NUM_THREADS; i++){
            thread_info[i].stars = &stars;
            thread_info[i].start = i*N/NUM_THREADS;
            thread_info[i].end = (i == NUM_THREADS-1) ? N : (i
            thread_info[i].t_G = G;
            thread_info[i].t_dt = dt;
            thread_info[i].t_N = N;

            pthread_create(&threads[i], NULL, update_stars, &t

        }
        for (int i = 0; i < NUM_THREADS; i++){
            pthread_join(threads[i], NULL);
        }
        count++;
    }
}
```

So now each thread struct reads all the arguments they need to create their section of the job.

3.3 OpenMP

The function `update_stars` is very similar to the serialized code, with a few differences:

```
/* Calculates and updates for new velocity and position of single star */
void update_stars(cbody_t *stars, const double G, const int N, const double dt){
    double *xp = stars->xp;
    double *yp = stars->yp;
    double *mass = stars->mass;
    double *xv = stars->xv;
    double *yv = stars->yv;
    const double Gdt = G*dt;

    /* Initialize loop vars here */
    double ix, iy, axi, ayi, jx, jy, mj, rx, ry, rsqrt, inv_r, inv_r_mj;
    int j, i;
    #pragma omp parallel for private(ix, iy, axi, ayi, j, jx, jy, mj, rx, ry)
    for (i = 0; i < N; i++) {
        ix = xp[i];
        iy = yp[i];
        axi = 0;
        ayi = 0;
        for (j = 0; j < N; j++) {
            jx = xp[j];
            jy = yp[j];
            mj = mass[j];
            rx = ix - jx;
            ry = iy - jy;
            rsqrt = sqrt(rx*rx + ry*ry) + eps;
            inv_r = 1.0/(rsqrt*rsqrt*rsqrt);
            inv_r_mj = inv_r*mj;
            axi += rx*inv_r_mj;
            ayi += ry*inv_r_mj;
        }
        xv[i] += -Gdt*axi;
        yv[i] += -Gdt*ayi;
    }
}
```

```
        /* Wait for threads to finish calculate all forces before updating posit  
        #pragma omp barrier  
        /* Static scheduling for better data locality */  
        #pragma omp parallel for schedule(static)  
        for (int i = 0; i < N; i++){  
            xp[i] += xv[i]*dt;  
            yp[i] += yv[i]*dt;  
        }  
    }
```

So in this case the calculations work similar except that like for pthreads the inner loop has to be changed as well as a new loop for updating positions had to be created, both to prevent race conditions. Static scheduling was used for both outer loops since we have better data locality when the star attributes comes as an array. There is no race conditions here since no thread will update any star simultaneously.

3.4 Performance and discussion

3.4.1 Results

The following are time measurements for 3000 and 5000 stars for 100 steps. Y axis is in seconds and X axis are number of maximum threads set. The timings were done using the ./time command in ubuntu, meaning the time to run the whole code is measured. Both user time and real time was noted. As previously stated the other input arguments are the same. All codes were compiled using -O3, and Pthreads with -pthread and OpenMP with -fopenmp.

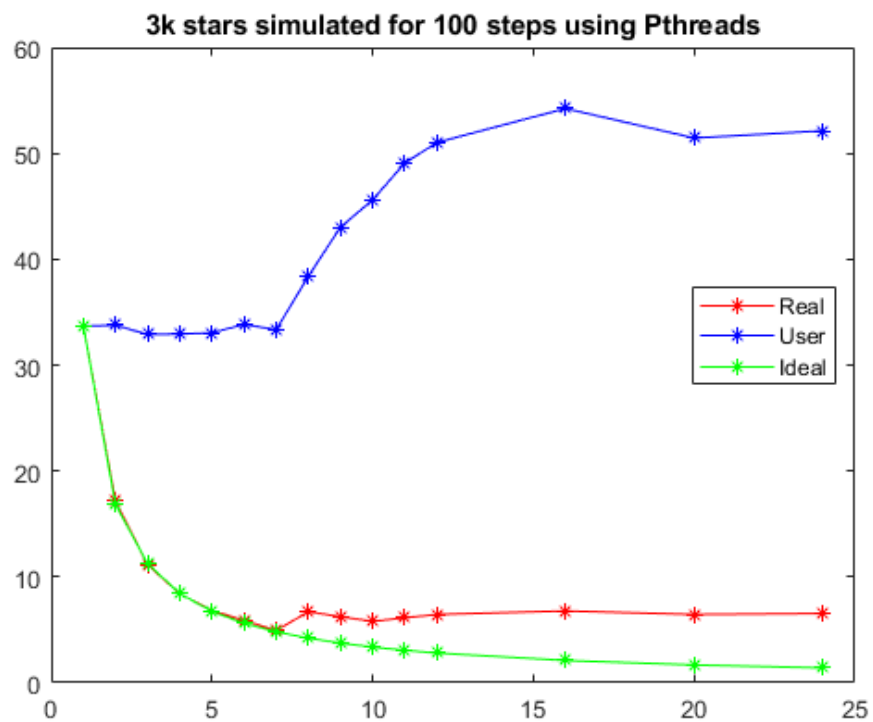


Figure 4: 3000 stars simulated and timed with Pthreads




Figure 5: 3000 stars simulated and timed with OpenMP

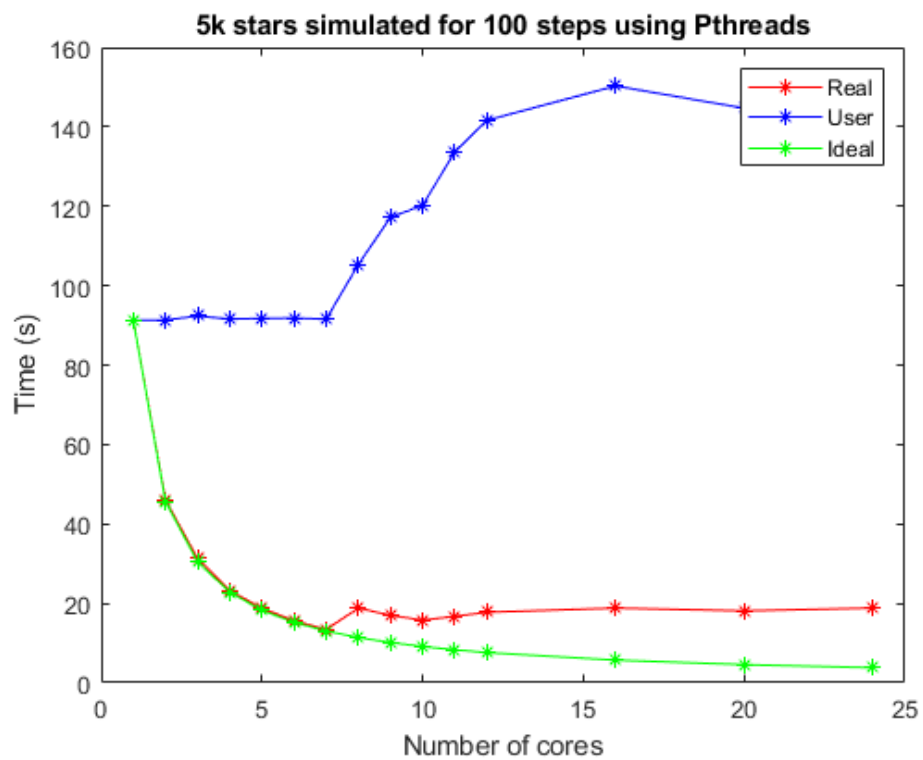


Figure 6: 5000 stars simulated and timed with Pthreads

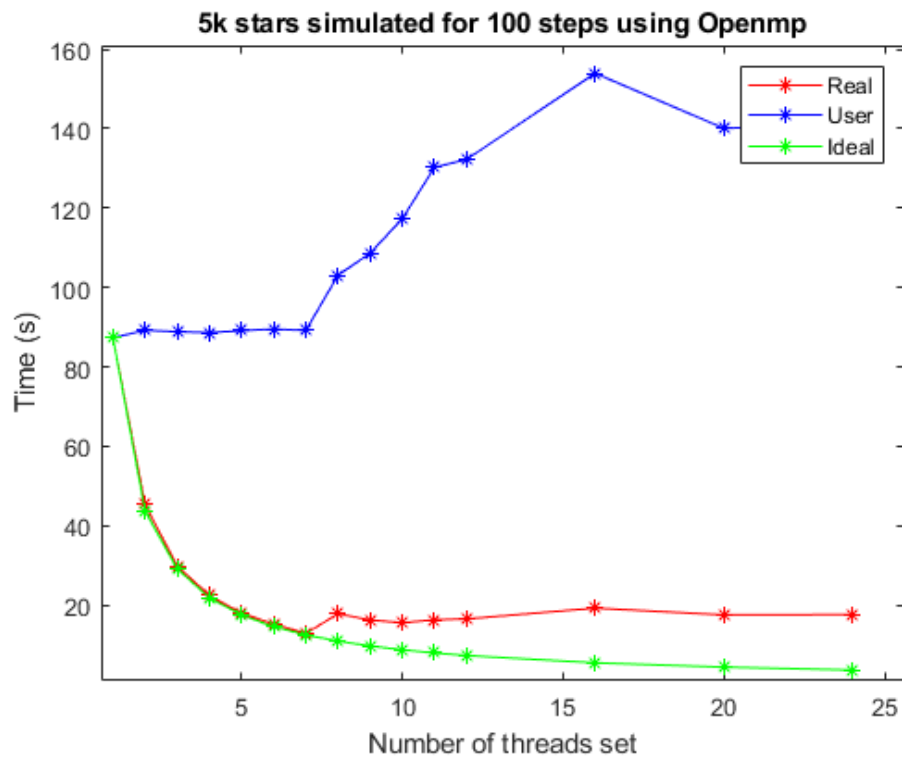


Figure 7: 5000 stars simulated and timed with OpenMP

To see how the difference between real time and ideal time grows for problemsize (here from 3000 stars to 5000 stars) we get the following plots:

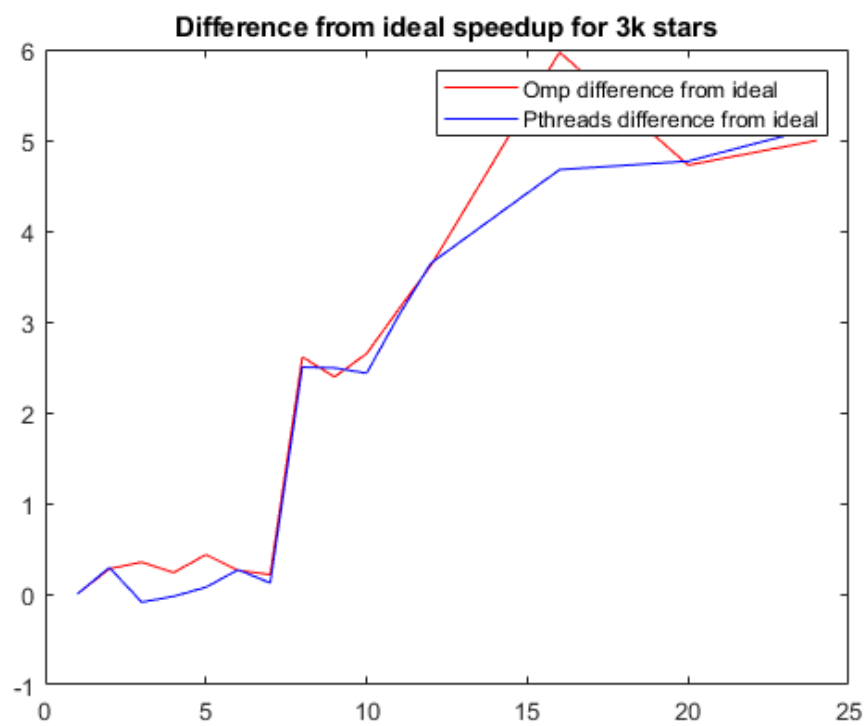


Figure 8: Real time - Ideal time for 3000 stars

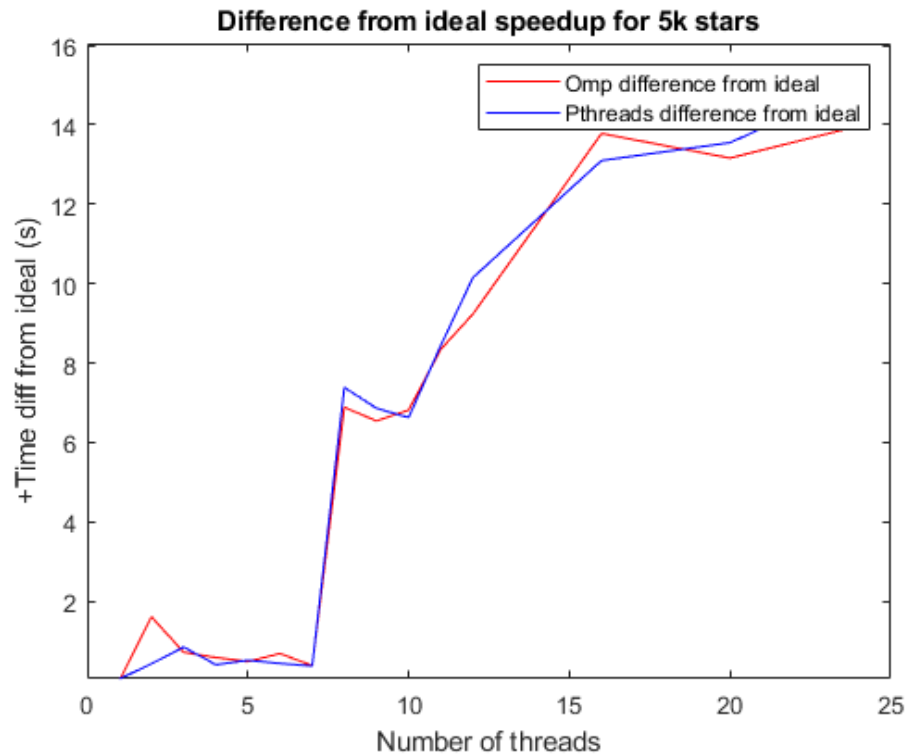


Figure 9: Real time - Ideal time for 5000 stars

Where it would seem that the difference scales linearly with problemsize, and there is no noticable improvement in the speedup for the bigger problem in my case.

The code still seems to follow the complexity of the serial code, where the y axis is time in seconds and the x axis is number of stars in the simulation:

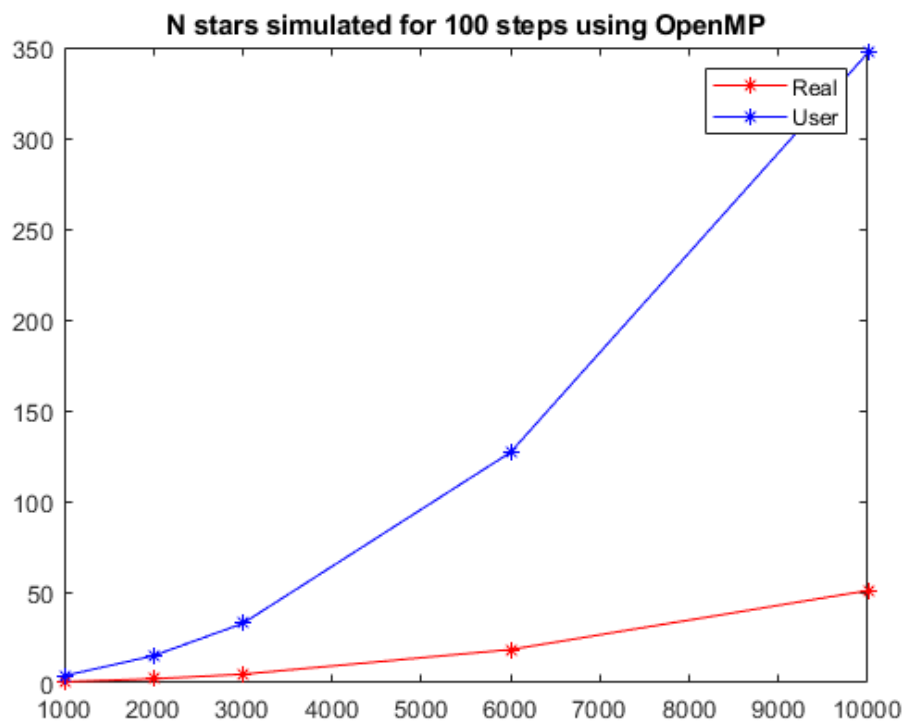


Figure 10: Time scaling with increased problem size using OpenMP

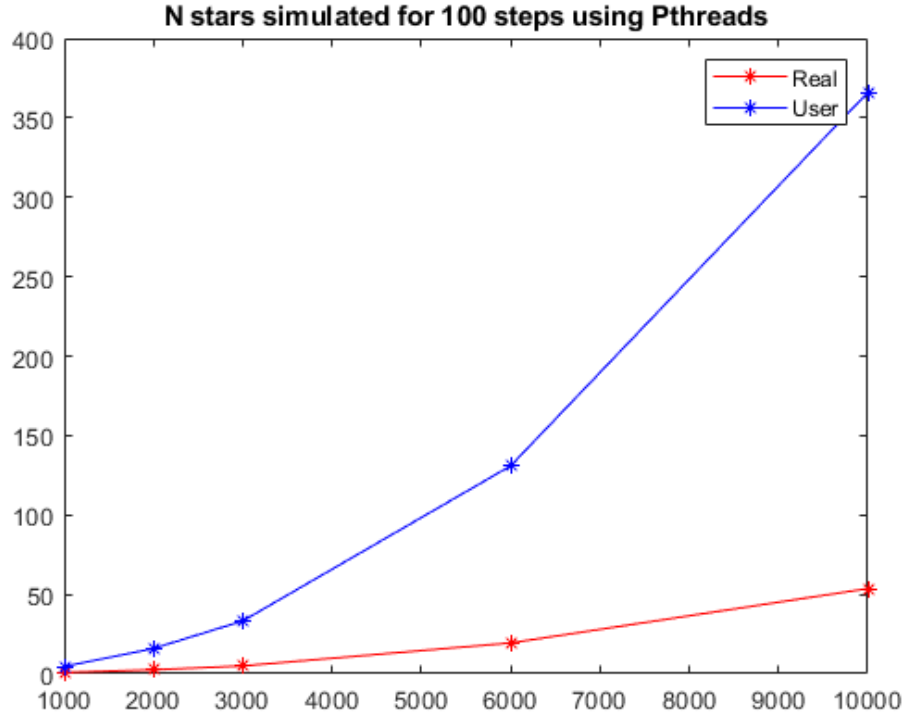


Figure 11: Time scaling with increased problem size using Pthreads

3.4.2 Discussion

In this part I used the same computer (fredholm) with the same hardware. There was problem parallelizing the loop as it was from the serial code. The inner loop had dependencies on other iterations and accessed the same memory at the same time, which made the calculations slow and inaccurate. The parallelized code runs slower on one core than the serialized code does (31.374 seconds for OpenMP, 33.710 seconds for Pthreads compared to 18.186 seconds for the serialized code) since I now have two outer loops in the function to avoid race conditions. Despite this, both Pthread and OpenMP implementation achieved much better results with the appropriate amount of cores selected. Compared to the serial code, OpenMP achieved a speed-up of 3.9 and Pthreads ≈ 3.7 . Compared to their respective baseline time (time for running on 1 thread), OpenMP achieved a speed-up of ≈ 6.68 and Pthreads ≈ 6.8 . The optimal number of threads to run on was 7 (CPU usage

on fredholm is capped at 800 percent)